

Assignment 2

Felix Söderman DD2360 HT23

Github: <https://github.com/felixcool200/DD2360HT23>

Exercise 1 - Reflection on GPU-accelerated Computing

1. List the main differences between GPUs and CPUs in terms of architecture.

ANSWER: The main difference between the two is that CPUs are based on Latency oriented processor architecture. This means that it is designed to minimize the time it takes to complete a single task. This results in a few cores that have high clock speed (and high energy usage). A CPU is thus good at solving single threaded tasks fast, such as sorting elements in a list.

GPUs on the other hand use a throughput-oriented processor architecture. This means it is designed to maximize the amount of problems it can solve rather than making sure they are completed as soon as they are created. GPUs thus have many cores and is good at completing task that are parallelizable, such as AI workloads or processing individual pixels on a screen.

2. Check the latest Top500 list that ranks the top 500 most powerful supercomputers in the world. In the top 10, how many supercomputers use GPUs Report the name of the supercomputers and their GPU vendor (Nvidia, AMD, ...) and model.

ANSWER: Since the definition of GPU is quite vague. I will also include computers that run other type of accelerator cards that have similar capabilities.

1 Frontier - HPE Cray EX235a, **AMD** Instinct MI250X

3 LUMI - HPE Cray EX235a, **AMD** Instinct MI250X

4 Leonardo - BullSequana XH2000, **NVIDIA** A100 SXM4 64 GB

5 Summit - IBM Power System AC922, **NVIDIA** Volta GV100

6 Sierra - IBM Power System AC922, **NVIDIA** Volta GV100

8 Perlmutter - HPE Cray EX235n, **NVIDIA** A100 SXM4 40 GB

9 Selene - NVIDIA DGX A100, **NVIDIA** A100

10 Tianhe-2A - TH-IVB-FEP Cluster, Intel Xeon E5-2692v2 12C 2.2GHz, **National University of Defense Technology (NUDT)** Matrix-2000 NOTE that this last one is not a GPU per say but an accelerator card.

3. One main advantage of GPU is its power efficiency, which can be quantified by Performance/Power, e.g., throughput as in FLOPS per watt power consumption. Calculate the power efficiency for the top 10 supercomputers. (Hint: use the table in the first lecture)

ANSWER: One can clearly see that the machines running GPUs are much more power efficient than those who does not.

- 1 Frontier - RMax = 1194 [PFlop/s], Power = 22703 [kW] => Power efficiency = 52.59 [Gflop/watts]
- 2 Supercomputer Fugaku - RMax = 442.01 [PFlops/s], Power = 29,899.23 kW [kW] => Power efficiency = 14.78 [Gflop/watts]
- 3 LUMI - RMax = 309.10 [PFlop], Power = 6015.77 [kW] => Power efficiency = 51.38 [Gflop/watts]
- 4 Leonardo - RMax = 238.70 [PFlop], Power = 7404.40 [kW] => Power efficiency = 32.24 [Gflop/watts]
- 5 Summit - RMax = 148.60 [PFlop], Power = 10096.00 [kW] => Power efficiency = 14.72 [Gflop/watts]
- 6 Sierra - RMax = 94.64 [PFlop], Power = 7438.28 [kW] => Power efficiency = 12.72 [Gflop/watts]
- 7 Sunway TaihuLight - RMax = 93.01 [PFlop], Power = 15371 [kW] => Power efficiency = 6.05 [Gflop/watts]
- 8 Perlmutter - RMax = 70.87 [PFlop], Power = 2589 [kW] => Power efficiency = 27.37 [Gflop/watts]
- 9 Selene - RMax = 63.46 [PFlop], Power = 2646 [kW] => Power efficiency = 23.98 [Gflop/watts]
- 10 Tianhe-2A - RMax = 61.44 [PFlop], Power = 18482 [kW] => Power efficiency = 3.32 [Gflop/watts]

Exercise 2 - Device Query

1. The screenshot of the output from you running deviceQuery test.

Screenshot from running Device Query on my laptop

```

./deviceQuery/deviceQuery Starting...

  CUDA Device Query (Runtime API) version (CUDART static linking)

Detected 1 CUDA Capable device(s)

Device 0: "NVIDIA GeForce GTX 1050 Ti with Max-Q Design"
  CUDA Driver Version / Runtime Version      12.2 / 12.2
  CUDA Capability Major/Minor version number: 6.1
  Total amount of global memory:             4041 MBytes (4237164544 bytes)
  (006) Multiprocessors, (128) CUDA Cores/MP: 768 CUDA Cores
  GPU Max Clock rate:                        1418 MHz (1.42 GHz)
  Memory Clock rate:                         3504 Mhz
  Memory Bus Width:                          128-bit
  L2 Cache Size:                             1048576 bytes
  Maximum Texture Dimension Size (x,y,z)     1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
  Total amount of constant memory:            65536 bytes
  Total amount of shared memory per block:    49152 bytes
  Total shared memory per multiprocessor:     98304 bytes
  Total number of registers available per block: 65536
  Warp size:                                 32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:        1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size   (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                     2147483647 bytes
  Texture alignment:                         512 bytes
  Concurrent copy and kernel execution:       Yes with 2 copy engine(s)
  Run time limit on kernels:                  Yes
  Integrated GPU sharing Host Memory:          No
  Support host page-locked memory mapping:     Yes
  Alignment requirement for Surfaces:          Yes
  Device has ECC support:                     Disabled
  Device supports Unified Addressing (UVA):     Yes
  Device supports Managed Memory:              Yes
  Device supports Compute Preemption:          Yes
  Supports Cooperative Kernel Launch:          Yes
  Supports MultiDevice Co-op Kernel Launch:    Yes
  Device PCI Domain ID / Bus ID / location ID: 0 / 1 / 0
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 12.2, CUDA Runtime Version = 12.2, NumDevs = 1
Result = PASS

```

Screenshot from running Device Query on Google Colab

```
./deviceQuery/deviceQuery Starting...

CUDA Device Query (Runtime API) version (CUDA static linking)

Detected 1 CUDA Capable device(s)

Device 0: "Tesla T4"
  CUDA Driver Version / Runtime Version      12.0 / 11.8
  CUDA Capability Major/Minor version number: 7.5
  Total amount of global memory:             15102 MBytes (15835398144 bytes)
  (040) Multiprocessors, (064) CUDA Cores/MP: 2560 CUDA Cores
  GPU Max Clock rate:                       1590 MHz (1.59 GHz)
  Memory Clock rate:                        5001 Mhz
  Memory Bus Width:                         256-bit
  L2 Cache Size:                           4194304 bytes
  Maximum Texture Dimension Size (x,y,z)    1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
  Total amount of constant memory:           65536 bytes
  Total amount of shared memory per block:   49152 bytes
  Total shared memory per multiprocessor:    65536 bytes
  Total number of registers available per block: 65536
  Warp size:                                32
  Maximum number of threads per multiprocessor: 1024
  Maximum number of threads per block:      1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                     2147483647 bytes
  Texture alignment:                        512 bytes
  Concurrent copy and kernel execution:     Yes with 3 copy engine(s)
  Run time limit on kernels:                 No
  Integrated GPU sharing Host Memory:        No
  Support host page-locked memory mapping:   Yes
  Alignment requirement for Surfaces:        Yes
  Device has ECC support:                    Enabled
  Device supports Unified Addressing (UVA):   Yes
  Device supports Managed Memory:            Yes
  Device supports Compute Preemption:        Yes
  Supports Cooperative Kernel Launch:        Yes
  Supports MultiDevice Co-op Kernel Launch:  Yes
  Device PCI Domain ID / Bus ID / location ID: 0 / 0 / 4
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 12.0, CUDA Runtime Version = 11.8, NumDevs = 1
Result = PASS
```

2. What architectural specifications do you find interesting and critical for performance? Please provide a brief description.

ANSWER: When looking at the specification I find that the:

GPU Max clock rate, amount of cores and streaming multiprocessors are important since they determine how fast each core computes and how many cores can work at the same time.

Memory clock rate, global memory size are also important since loading in assets from the CPU is much slower than from GDRAM. The cache size can also improve performance since it is even faster than GDRAM.

These are the most relevant and critical features for performance.

3. How do you calculate the GPU memory bandwidth (in GB/s) using the output from deviceQuery? (Hint: memory bandwidth is typically determined by clock rate and bus width, and check what double data rate (DDR) may impact the bandwidth)

ANSWER: To calculate the memory bus speed GB/s. First translate the memory bus width from bits to Bytes by dividing it by 8. Then take the Memory Clock rate and double it to get the double data

rate(DDR). After that take the new bus width [B] divided by the new memory clock lock speed [Hz] which then equals $[B \cdot \text{Hz}] = [B \cdot (1/\text{s})] = [\text{B}/\text{s}]$. Lastly divide it 10^9 to make it [GB/s].

Thus in my case:

```
Bus Width:
128 bit/8 = 16 Bytes.

Clock Rate:
3504 MHz*2 = 7008 MHz.

Memory bus speed = (16 * (7008 * 10^6)) / 10^9 = 112.128 GB/s.
```

4. Compare your calculated GPU memory bandwidth with Nvidia published specification on that architecture. Are they consistent?

ANSWER: When looking up the speed I get 112.1 GB/s which is identical to what was calculated.

Exercise 3 - Compare GPU Architecture

1. List 3 main changes in architecture (e.g., L2 cache size, cores, memory, notable new hardware features, etc.)

ANSWER: I decided to look at the consumer grade GPUs and there the three latest GPU architecture are Ada Lovelace (4000-series), Ampere (3000-series) and Turing (2000 -series). I decided to test the three best customer grade cards from each generation.

Card	RTX 2080 Ti RTX	RTX 3090 Ti RTX	GeForce RTX 4090
L2 Cache (MB)	5.5 MB	6 MB	72 MB
Tensor Cores	544	336	512
SM Count	68	84	128
SP Count (Per SM)	64	128	128
NVIDIA CUDA-cores	4352	10752	16384
Boost clock (GHz)	1.64 GHz	1.86 GHz	2.52 GHz
Base clock (GHz)	1.35 GHz	1.56 GHz	2.23 GHz
Bus Interface	PCIe 3.0 x16	PCIe 4.0 x16	PCIe 4.0 x16
Global memory (GB)	11 GB	24 GB	24 GB
Memory type	GDDR6	GDDR6X	GDDR6X

2. List the number of SMs, the number of cores per SM, the clock frequency and calculate their theoretical peak throughput.

ANSWER: Throughput is calculated as $SMs * SPs * Clock\ (GHz) * 1/1000 = Troughput\ (Tflops)$

Card	RTX 2080 Ti RTX	RTX 3090 Ti RTX	GeForce RTX 4090
SM Count	68	84	128
SP Count (Per SM)	64	128	128
Boost clock (GHz)	1.64 GHz	1.86 GHz	2.52 GHz
Max Troughput (Tflops)	7.14	20.00	41.29

3. Compare (1) and (2) with the NVIDIA GPU that you are using for the course.

ANSWER: I run three different GPUs in the course (1080 ti, 1050 ti Max-Q and Google colab (T4)). I will compare them with a GTX 1080 Ti

Card	RTX 1080 Ti RTX
L2 Cache (MB)	2.75 MB
Tensor Cores	0
SM Count	28
SP Count (Per SM)	128
NVIDIA CUDA-cores	3584
Boost clock (GHz)	1.58 GHz
Base clock (GHz)	1.428 GHz
Bus Interface	PCIe 3.0 x16
Global memory (GB)	11 GB
Memory type	GDDR5X
Max Troughput (Tflops)	5.66

Exercise 4 - Rodinia CUDA benchmarks and Profiling

1. Compile both OMP and CUDA versions of your selected benchmarks. Do you need to make any changes in Makefile?

ANSWER: When running the CUDA version of particalfilter the compute capability had to be changed to match my GPU by changing `sm_13` to `sm_61` on line 12 and line 15 in the makefile.

When running lavaMD no changes was needed.

When running hotspot3D the first line in the makefile as incorrect and needed to be changed from

`include ~/rodinia_3.0/common/make.config`
to

```
include ../../common/make.config
```

2. Ensure the same input problem is used for OMP and CUDA versions. Report and compare their execution time.

ANSWER:

Running particalfilter

OpenMP

```
make clean
make openmp

cat run
./particle_filter -x 128 -y 128 -z 10 -np 100000
```

RESULTS:

```
./run
ENTIRE PROGRAM TOOK 5.897769
```

CUDA

```
make clean
make all

cat run
./particlefilter_naive -x 128 -y 128 -z 10 -np 100000
```

RESULTS:

```
./run
ENTIRE PROGRAM TOOK 1.680555
```

Running lavaMD

OpenMP

```
make clean
```

```
cat run
./lavaMD -cores 12 -boxes1d 35
```

RESULTS:

```
./run
Configuration used: cores = 12, boxes1d = 35
Time spent in different stages of CPU/MCPU KERNEL:
0.000000000000 s, 0.000000000000 % : CPU/MCPU: VARIABLES
0.000006000000 s, 0.000024581621 % : MCPU: SET DEVICE
0.000000000000 s, 0.000000000000 % : CPU/MCPU: INPUTS
24.408475875854 s, 99.999984741211 % : CPU/MCPU: KERNEL
Total time:
24.408479690552 s
```

CUDA

```
make clean

cat run
./lavaMD -boxes1d 35
```

RESULTS:

```
./run
thread block size of kernel = 128
Configuration used: boxes1d = 35
Time spent in different stages of GPU_CUDA KERNEL:
0.064755000174 s, 0.542295157909 % : GPU: SET DEVICE / DRIVER INIT
0.000539999979 s, 0.004522267263 % : GPU MEM: ALO
0.034356001765 s, 0.287716686726 % : GPU MEM: COPY IN
11.818737983704 s, 98.976829528809 % : GPU: KERNEL
0.014343000017 s, 0.120116434991 % : GPU MEM: COPY OUT
0.008182000369 s, 0.068520717323 % : GPU MEM: FRE
Total time:
11.940914154053 s
```

Running hotspot3D**OpenMP:**

```
make clean
make 3D
```



```
cat run
./3D 512 8 10000 ../../data/hotspot3D/power_512x8
../../data/hotspot3D/temp_512x8 output.out

./run
```

RESULTS:

```
12 threads running
Time: 38.112 (s)
Accuracy: 4.856862e-05
```

CUDA

```
make clean
make release

cat run
./3D 512 8 10000 ../../data/hotspot3D/power_512x8
../../data/hotspot3D/temp_512x8 output.out

./run
```

RESULTS:

```
Time: 9.737 (s)
Accuracy: 4.096975e-05
```

3. Do you observe expected speedup on GPU compared to CPU? Why or Why not? **ANSWER:** Over all the three workloads the CUDA version ran faster. This is most likely since all the workload were easily parallelizable. When running the programs with fewer iterations/smaller values (for example amount of particles in the particlefilter) the CPU and GPU had comparable speeds but when increasing the amount of particles in parallel the CUDA program ran much faster.

(Bonus) Exercise 5 - GPU Architecture Limitations and New Development

ANSWER: I decided to read "Buddy compression: Enabling larger memory for deep learning and HPC workloads on GPUs" [1]:

1. What limitations this paper proposes to address?

ANSWER: This paper tries to address the problem of (relatively) low high speed memory on GPUs. Which can limits the GPU core to be fully utilized if each thread needs to read much data from the global memory.

2. What workloads/applications does it target?

ANSWER: The paper targets deep learning and HPC workloads since they are sterotypical workloads that have high memory usage, but any workload which requires much data to be read is applicable.

3. What new architectural changes does it propose? Why it can address the targeted limitation?

ANSWER: They propose of this paper is to evaluate and check what memory compression algorithms would work as a memory-expansion alternative. And to be used as a slower-but-larger buddy memory connected thorough a high-bandwidth interconnect, such as NVLink.

4. What applications are evaluated and how did they setup the evolution environment (e.g., simulators, real hardware, etc)?

ANSWER: The majority of their testing was using 16 different HPC/Deep learning benchmarks that where tested using multiple algortims and while messuring many different things. Thus mostly simulators where used. (They did compare UM and buddy compression on real hardware in one test).

5. Do you have any doubts or comments on their proposed solutions?

ANSWER: I myself does not any any doubts, but I am sceptical to how much of an improvment this will be. As it seems to add a lot of complexity whist giving a quite small gain.

References

[1] Choukse, E., Sullivan, M. B., O'Connor, M., Erez, M., Pool, J., Nellans, D., & Keckler, S. W. (2020, May). Buddy compression: Enabling larger memory for deep learning and HPC workloads on GPUs. In 2020 ACM/IEEE 47th Annual International Symposium on Computer Architecture (ISCA) (pp. 926-939). IEEE.